

Complex Data Types



Paul O'Fallon

@paulofallon

Complex Data Types

Collections

Tuples

User Defined
Types

Collections: Set

More info

Level **Intermediate**


Rating ★★★★★

Duration **2h 48m**

Released **19 Dec 2012**

Features 

```
CREATE TABLE courses (  
    id varchar,  
    name static,  
    // ...  
    features set<varchar> static,  
    module_id int,  
    // ...  
    PRIMARY KEY (id, module_id)  
);
```



Collections: Set

Inserting with a set

```
INSERT INTO courses (id, name, features)
VALUES ('node-intro', 'Introduction to Node.js', {'cc'});
```

Adding to a set

```
UPDATE courses SET features = features + {'cc'}
WHERE course_id = 'node-intro';
```

Collections: Set

Removing from a set

```
UPDATE courses SET features = features - {'cc'}  
WHERE course_id = 'node-intro';
```











Emptying the entire set

```
UPDATE courses SET features = {}  
WHERE course_id = 'node-intro';
```

Collections: List

Table of contents [Expand all](#)

0%

 Introduction and History		19:15
Introduction		1:01
Container History		4:56
Advantages		1:47
Advantages Broken Out		2:52
Docker, The Container Story		3:26
Docker Community		3:22
Docker Technology		1:34
Summary		0:17

Collections: List

```
CREATE TABLE courses (  
    id varchar,  
    name static,  
    // ...  
    module_id int,  
    clips list<varchar>,  
    // ...  
    PRIMARY KEY (id, module_id)  
);
```

Collections: List

Inserting with a list

```
INSERT INTO courses (id, module_id, clips)
VALUES ('docker-fundamentals', 1, ['Container History']);
```

Adding to a list

```
UPDATE courses SET clips = ['Introduction'] + clips
WHERE course_id = 'docker-fundamentals' AND module_id = 1;
```

```
UPDATE courses SET clips = clips + ['Advantages']
WHERE course_id = 'docker-fundamentals' AND module_id = 1;
```


Collections: List

Removing from a list


```
UPDATE courses SET clips = clips - ['Introduction']  
WHERE course_id = 'docker-fundamentals' and module_id = 1;
```

Manipulating a list by element id

```
UPDATE courses SET clips[3] = 'Advantages Broken Out'  
WHERE course_id = 'docker-fundamentals' AND module_id = 1;
```

```
DELETE clips[3] FROM courses  
WHERE course_id = 'docker-fundamentals' AND module_id = 1;
```

Collections: Map



Collections: Map

```
CREATE TABLE users (  
    id varchar,  
    first_name varchar,  
    last_name varchar,  
    password varchar,  
    reset_token varchar,  
    last_login map<varchar,timestamp>,  
    PRIMARY KEY (id)  
);
```

Collections: Map

Inserting with a map

```
INSERT INTO users (id, first_name, last_name, last_login)
VALUES ('john-doe', 'John', 'Doe',
        { '383cc0867cd2': '2015-06-30 09:02:24' });
```

Updating / adding to a map

```
UPDATE users SET last_login['383cc0867cd2']
    = '2015-07-01 11:17:42' WHERE user_id = 'john-doe';
```

```
UPDATE users SET last_login = last_login + { '7eb0a8997f39':
    '2015-07-02 07:32:17' } WHERE user_id = 'john-doe';
```

Collections: Map

Removing from a map

```
DELETE last_login['383cc0867cd2'] FROM users  
WHERE id = 'john-doe';
```

```
UPDATE users SET last_login = last_login - {'7eb0a8997f39'}  
WHERE id = 'john-doe';
```

Emptying the entire map

```
UPDATE users SET last_login = {}  
WHERE id = 'john-doe';
```

Collections and TTL

```
UPDATE users USING TTL 31536000  
SET last_login['383cc0867cd2'] = '2015-07-01 11:17:42'  
WHERE user_id = 'john-doe';
```

Tuples

(varchar, int, int, varchar, timestamp)

Tuples

383cc0867cd2



2015-07-01 11:17:42



383cc0867cd2



2015-07-01 11:17:42



98.203.153.64

Tuples

```
CREATE TABLE users (  
    id varchar,  
    first_name varchar,  
    last_name varchar,  
    password varchar,  
    reset_token varchar,  
    last_login map<varchar,  
        frozen<tuple<timestamp, inet>>>,  
    PRIMARY KEY (id)  
);
```











"Frozen"

- Nested types are serialized as a single (blob) value
- True for nested collections as well (`list<set<varchar>>`)
- Nested values must be set or read as a whole
- `frozen<>` makes this distinction obvious
- Leaves open the possibility of "non-frozen" support in the future

User Defined Types

Table of contents [Expand all](#)

0%

 Introduction and History		19:15
Introduction		1:01
Container History		4:56
Advantages		1:47
Advantages Broken Out		2:52
Docker, The Container Story		3:26
Docker Community		3:22
Docker Technology		1:34
Summary		0:17

User Defined Types

```
CREATE TYPE clip (name varchar, duration int);
```

```
CREATE TABLE courses (  
    id varchar,  
    author varchar static,  
    // ...  
    clips list<frozen<clip>>,  
    module_id int,  
    // ...  
    PRIMARY KEY (id, module_id)  
);
```

Why Not Just Use a Tuple?

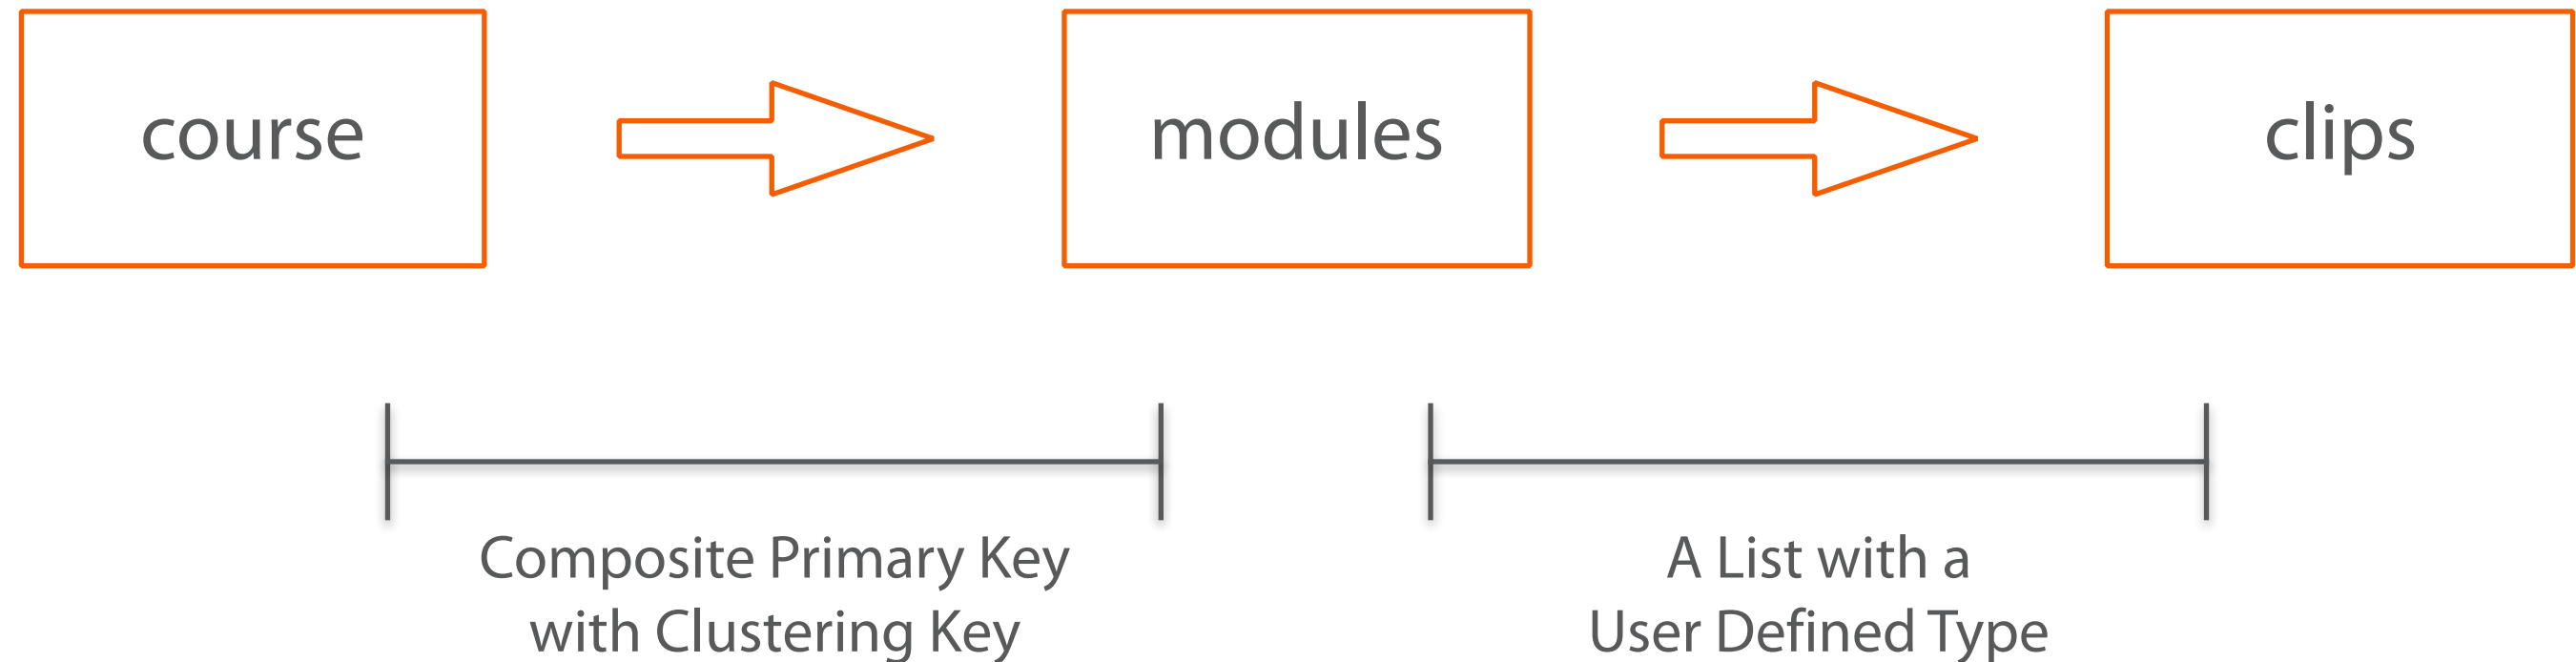
- Can identify individual components by name (not just order)
- Even more helpful with multiple components of the same type
- Helpful to start with a Tuple when modeling a User Defined Type
- Opportunity to reuse User Defined Type across tables

User Defined Types

```
CREATE TYPE person (name varchar, id varchar);
```

```
CREATE TABLE courses (  
    id varchar,  
    author frozen<person> static,  
    // ...  
    clips list<frozen<clip>>,  
    module_id int,  
    // ...  
    PRIMARY KEY (id, module_id)  
);
```

Courses, Modules and Clips



All in a single partition!

Conclusion

- Sets, Lists and Maps
- Collections and TTLs
- Tuples
- “Frozen” Nested Complex Types
- User Defined Types