

Making the Most of Cassandra



Paul O'Fallon

@paulofallon

Making the Most of Cassandra

Indexes

Batches

Transactions

Secondary Indexes

```
CREATE TABLE users (  
    id varchar,  
    first_name varchar,  
    last_name varchar,  
    company varchar,  
    // ...  
    PRIMARY KEY (id)  
);
```

```
CREATE INDEX users_company ON users (company);
```

Secondary Indexes

```
INSERT INTO users (id, first_name, last_name, company)  
VALUES ('john-doe', 'John', 'Doe', 'acme-corp');
```

```
SELECT * FROM users WHERE company = 'acme-corp';
```

Secondary Indexes on Collections

Tags

developer (1334)

open-source (183)

javascript (174)

node.js (35)

```
CREATE TABLE users (  
    id varchar,  
    first_name varchar,  
    last_name varchar,  
    company varchar,  
    tags set<varchar>,  
    // ...  
    PRIMARY KEY (id)  
);
```

```
CREATE INDEX ON users (tags);
```

Secondary Indexes on Collections

```
INSERT INTO users (id, first_name, last_name, company, tags)  
VALUES ('john-doe', 'John', 'Doe', 'acme-corp', {'java'});
```

```
SELECT * FROM users WHERE tags CONTAINS 'java';
```

Secondary Indexes on Collections

✓ Set
✓ List
✓ Map

CREATE INDEX ON <table> (<collection_column>)

WHERE <collection_column> **CONTAINS** <value>



✓ Map

CREATE INDEX ON <table> (**KEYS** (<map_column>))

WHERE <map_column> **CONTAINS KEY** <key>



When Not to Use a Secondary Index

- High (or very low) cardinality columns
- Tables with counter columns
- Frequently updated or deleted columns
- Tables with very large partitions
- Static columns (for now)

Manually Maintained Indexes

Tags

developer (1334)

open-source (183)

javascript (174)

node.js (35)

```
CREATE TABLE courses (  
    id varchar,  
    // ...  
    tags set<varchar> static,  
    module_id int,  
    // ...  
    PRIMARY KEY (id, module_id)  
);
```

Manually Maintained Indexes

```
CREATE TABLE course_tags (  
    tag varchar,  
    course_id varchar,  
    PRIMARY KEY (tag, course_id)  
);
```

Manually Maintained Indexes

```
INSERT INTO courses (id, name, tags)
VALUES ('node-intro', 'Introduction to Node.js',
{'developer', 'open-source', 'javascript', 'node.js'});
```

```
INSERT INTO course_tags (tag, course_id)
VALUES ('developer', 'node-intro');
```

```
INSERT INTO course_tags (tag, course_id)
VALUES ('open-source', 'node-intro');
```

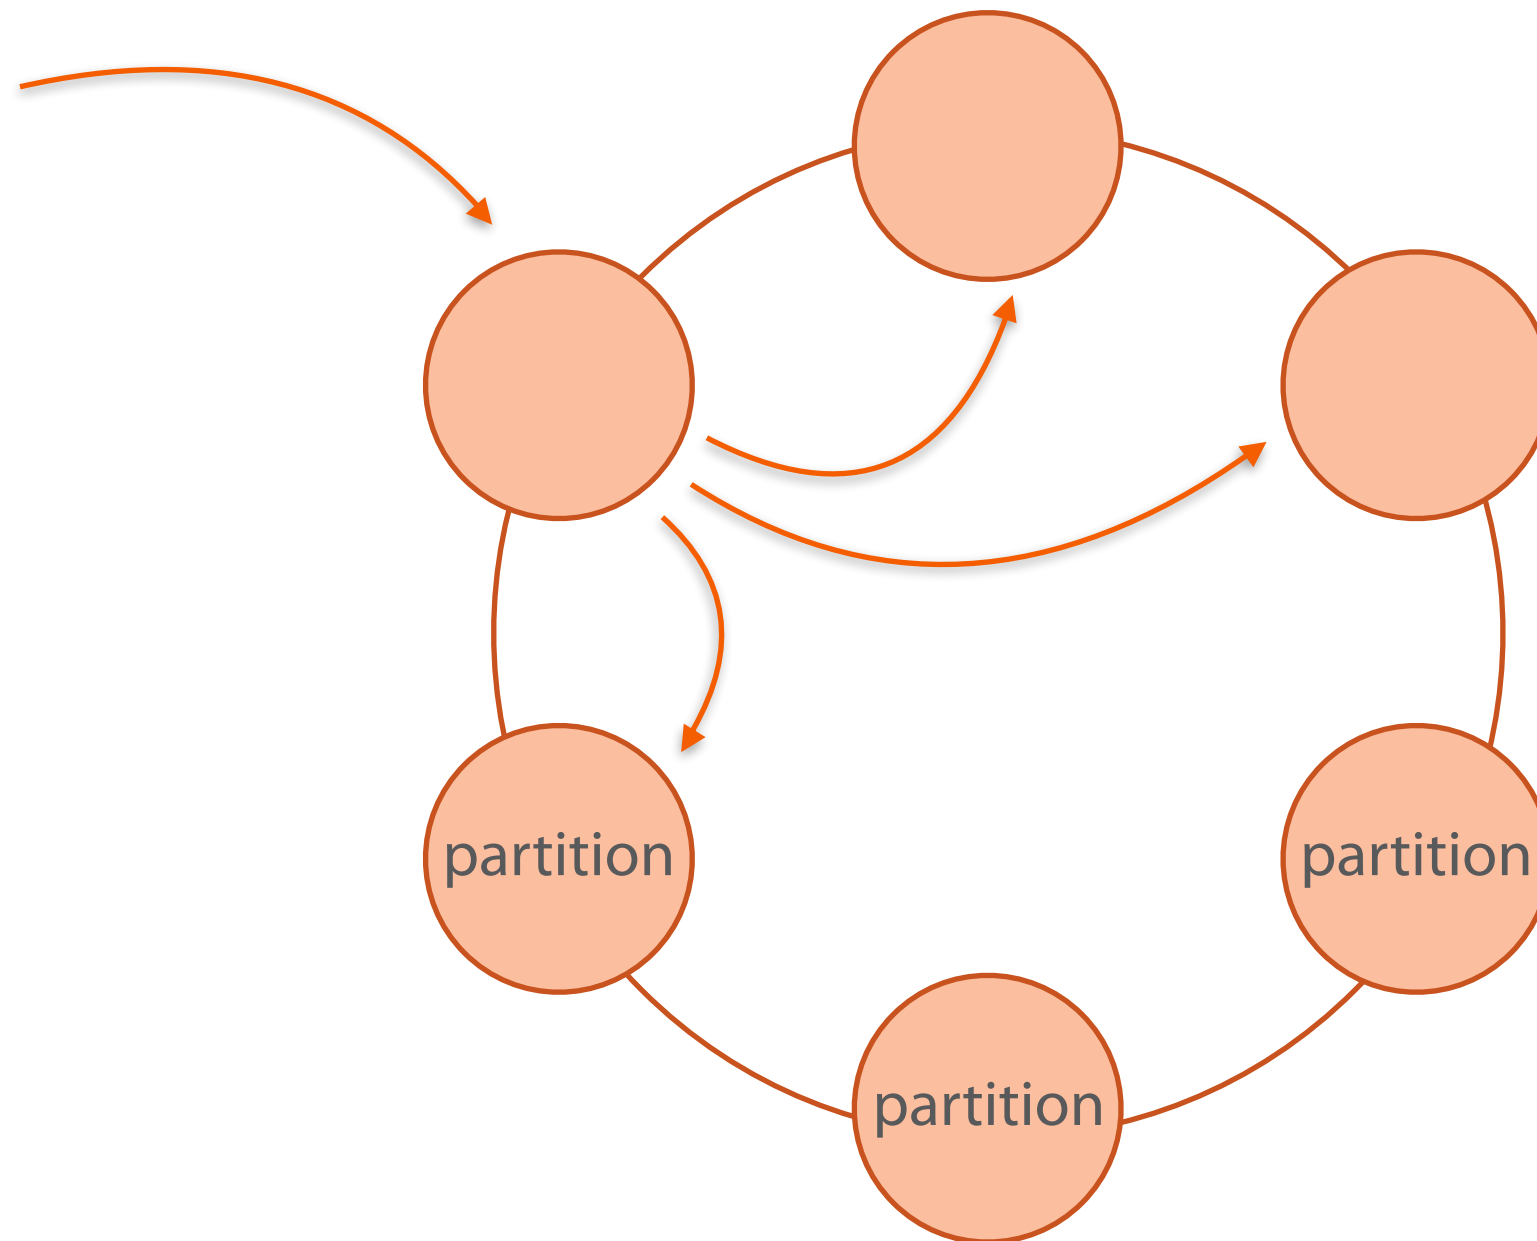
```
// ... etc.
```

Batches


Intended for:
Keeping tables in sync

Not intended for:
Fast loading of data

Batches



Batches



Batches

```
BEGIN BATCH
```

```
INSERT INTO courses (id, tags)  
VALUES ('node-intro', {'developer', 'javascript', 'node.js',  
  'open-source'});
```


```
INSERT INTO course_tags (tag, course_id, course_name)  
VALUES ('developer', 'node-intro', 'Introduction to Node.js');
```

```
// ... etc.
```

```
APPLY BATCH;
```

Unlogged Batches

**ONE
WRITE!**



Unlogged Batches

```
BEGIN UNLOGGED BATCH
```

```
INSERT INTO courses (id, name)  
VALUES ('node-intro', 'Introduction to Node.js');
```

```
INSERT INTO courses (id, module_id, module_name)  
VALUES ('node-intro', 1, 'Getting Started with Node.js');
```

```
// ... etc.
```

```
APPLY BATCH;
```


Alternative to Batches

Client pseudocode

```
PREPARE <insert_statement>  
LOOP_OVER_DATA  
    EXECUTE (<prepared_statement>, <row_of_data>)  
END_LOOP
```

Let the client optimize for you!

Lightweight Transactions



1. Prepare \leftrightarrow Promise
2. Read \leftrightarrow Results
3. Propose \leftrightarrow Accept
4. Commit \leftrightarrow Ack

Compare-and-Set Operations

Insert

```
INSERT INTO users (id, first_name, last_name)  
VALUES ('john-doe', 'John', 'Doe')  
IF NOT EXISTS;
```

Update

```
UPDATE users SET password = 'mypass', reset_token = null  
WHERE id = 'john-doe'  
IF reset_token = '1GRhEs1';
```

Works with Batches Too!

```
BEGIN BATCH
```

```
    INSERT INTO <table> ... IF NOT EXISTS;
```

```
    INSERT INTO <table> ...
```

```
    INSERT INTO <table> ...
```

```
APPLY BATCH;
```

Conclusion

- Secondary Indexes
- Indexed Collections
- Manually Maintained Indexes
- Batches (Logged & Unlogged)
- Lightweight Transactions